



DOL ERISA EFAST2 Modernization Vendor Industry Day

May 20, 2026



U.S. Small Business
Administration

[Satsyil.com](https://satsyil.com)

Agenda

1. Company Overview
2. EFAST2 Understanding
3. A. Proposed Technical Architecture (Kiran)
4. B. Security and Vulnerability Management (Kiran)
5. C. Data Quality and Validation (Josh)
6. D. Annual Rollover and Version Changes (Josh)
7. E. AI Use in Development, PII Identification, and Testing (Pierre)
8. F. Data Migration and Transition (Josh)
9. G. Third-Party and Bulk Filing (Kiran)
10. H. Contact Center (Don)
11. Similar Application Demonstration and Fallback Video (Kiran)
12. Government Areas of Interest (Kiran and team)
13. Q&A

Company Overview

Prime Contractor

Satsyil Corp

Satsyil Corp (Satsyil) is a CMMI Level 5, ISO 9001/27001 certified small business with 13+ years delivering mission-critical federal IT systems. We are a trusted **SBA-certified** innovator and trusted federal partner with **excellent CPARS** for high-volume federal filing systems, data modernization, and cloud-native DevSecOps. Contract Vehicle: GSA MAS (47QTCA18D009K)



Artificial Intelligence/ Machine Learning

AI/ML-driven automation and PII detection for high-volume federal filing systems (USPTO).



Cloud Engineering & Modernization

Delivered platforms processing 20M+ transactions/day across 10+ petabytes of federal data.



DevSecOps / Agile Transformation

Deployed 100+ DevSecOps pipelines with FedRAMP, NIST 800-53, and continuous ATO support.



Enterprises Data Services & Governance

Engineered enterprise data platforms supporting NARA's 300M+ digital records catalog.



Enterprise Custom App Dev

Proven CMMI Level 5 processes, Exceptional CPARS ratings, and ISO 9001/27001 certifications.

EFAST2 Understanding

Key Requirements

Platform Scope

- Process 1M+ annual Form 5500 filings across all schedules
- Web portal (IFILE) for direct filer submissions
- Bulk API intake for third-party software providers
- Participant Acceptance Testing System (PATS)
- Public data search and dissemination portal
- Contact center supporting filers and software developers
- Secure distribution to DOL, IRS, PBGC, and NARA
- Annual form rollover and schema update management

Key DOL Requirements

- 99.99% availability with auto-scaling cloud infrastructure
- FedRAMP High/Moderate authorized hosting environment
- Zero Trust security with NIST SP 800-53 Rev. 5 controls
- Login.gov integration at IAL-1/IAL-2 and AAL2
- Pre-submission validation and real-time error feedback
- Configurable business rules with no-code annual updates
- Full audit trail and submission lifecycle traceability
- NACI-cleared staff with 30-day transition-in capability

A. Proposed Technical Architecture

A1. Form Intake, Processing & Distribution

Two Intake Channels

- IFILE Web Portal: browser-based submission UI for direct filers; WCAG 2.2 AA; Login.gov authentication at IAL-1/IAL-2
- Bulk Filing API: RESTful JSON-first with backward-compatible XML for third-party software providers; API key authentication
- All submissions enter a controlled SQS intake layer — decoupled from processing, durable, and loss-proof even under peak load
- Unique tracking ID issued immediately upon intake — filer receives acknowledgment before processing begins
- Pre-validation executes before formal acceptance: schema check, field-level rules, entity validation (EIN, plan number)
- AWS Step Functions orchestrate the processing pipeline: validate → process → distribute — exactly-once, auditable, stateless
- Validated filings distributed via encrypted SFTP/API to DOL, IRS, and PBGC per the IRD; NARA annual archive packages prepared in non-proprietary formats with full metadata
- Public data dissemination via OpenSearch-powered search portal — processed filings made accessible per disclosure requirements

Processing Pipeline (5 Stages)

- ① Intake: Submission received via IFILE or Bulk API; SQS queued; tracking ID issued; acknowledgment sent
- ② Validate: Schema validated against DER; field-level, cross-field, and entity rules applied; errors returned with actionable guidance
- ③ Process: Step Functions execute processing workflow; exactly-once delivery enforced via idempotency keys
- ④ Acknowledge: Real-time status updates delivered to filer via API and dashboard at each milestone
- ⑤ Distribute: Encrypted delivery to DOL, IRS, PBGC drop-boxes; NARA archive; OpenSearch public portal updated
- All data in motion: TLS 1.3 + FIPS-validated cryptography
- All data at rest: AWS KMS customer-managed keys
- Every stage logged with timestamp and audit record — full DOL visibility into pipeline at any moment

A2. Application Components & Data Flows

Core Application Components

- IFILE Web Portal: public-facing filing UI; WCAG 2.2 AA; Login.gov SSO; responsive design
- Bulk Filing API Service: RESTful JSON/XML intake; API key auth; rate limiting; developer portal
- Intake & Queue Layer: Amazon SQS + Step Functions; exactly-once delivery; DLQ exception handling
- Validation Engine: schema, field-level, cross-field, entity validation; driven by versioned DER metadata; same engine in PATS and production
- Processing Service: stateless microservices on EKS; auto-scaling; parallel filing execution
- Status & Acknowledgment Service: real-time tracking API; filer dashboard; webhook push updates
- Admin Configuration Panel: DOL-accessible no-code UI for schema and rule updates; government approval workflow
- Distribution Service: encrypted SFTP/API to DOL, IRS, PBGC; OpenSearch public portal; NARA archive

Data Flows & Supporting Services

- Data flow direction: Intake → Validate → Process → Acknowledge → Distribute (single-direction, auditable at every stage)
- PATS Environment: production-mirrored sandbox for third-party certification; same validation engine and APIs
- Contact Center (AWS Connect + Lex IVR): CRM records tied to specific submission IDs; full call recording
- Amazon RDS (PostgreSQL): primary + read replicas; automated failover; FTI data in isolated encrypted repository
- ElastiCache (Redis): session management and validation result caching for high-throughput periods
- CloudWatch + PagerDuty: real-time monitoring; ops team alerted within 60 seconds of any anomaly
- AWS WAF + Shield: DDoS protection; CloudFront CDN offloads 90% of static traffic from origin
- All inter-service communication via private VPC endpoints — no public internet traversal for internal data flows

A3. Integration Approach

Five Integration Partners

- Login.gov via OIDC/SAML 2.0: IAL-1/IAL-2 and AAL2 for all public filers; PIV/HSPD-12 for government users; short-lived tokens and continuous session monitoring
- DOL, IRS & PBGC: validated filings transmitted via encrypted SFTP/API drop-box delivery per the IRD; TLS 1.3; retry logic ensures delivery even during downstream outages
- DOL Enterprise Services: Jira, SharePoint, and GitLab connected via VPC peering and private endpoints — no public internet traversal for internal data flows
- Third-Party Software Providers: RESTful JSON-first API with backward-compatible XML; developer portal with sandbox, docs, and code samples; PATS certification before production access
- NARA — Archival Delivery: annual archive packages in non-proprietary formats with structured metadata; hash-based integrity verification; FTI excluded per IRS Pub. 1075
- Contact Center (AWS Connect): CRM records tied to specific submission IDs; call recording and transcript search; real-time agent dashboards
- OpenSearch Public Portal: validated filings made searchable per public disclosure requirements; high-performance query API for researchers and third parties

Integration Design Principles

- API-first, standards-based: all integrations use documented, versioned APIs — no proprietary connectors or undocumented dependencies
- Version-controlled with 90-day deprecation windows: external partners are never broken by an update; backward-compatible XML supported alongside JSON
- Zero public internet for internal flows: all DOL enterprise integrations routed through VPC private endpoints and peering
- Authenticated at every boundary: Login.gov for filers, PIV/HSPD-12 for government users, scoped API keys with NIST 800-63 lifecycle management for third parties
- Retry and confirmation at every integration point: delivery to IRS, PBGC, and DOL confirmed with receipts; automated retry with exponential backoff on failure
- Integration monitoring: CloudWatch tracks every outbound delivery; failed deliveries alert ops team within 60 seconds with structured metadata
- Tested before cutover: all IRS, PBGC, and DOL drop-box integrations fully tested and certified in staging before any production migration gate
- Exit-ready: all integration configurations and credentials exportable;

A4. Hosting & Deployment

AWS GovCloud Hosting Environment

- AWS GovCloud (FedRAMP Moderate) — Satsyil's proven environment at USPTO and NARA; authority to operate on day one
- Amazon EKS (Kubernetes): stateless microservices orchestrated across 3 Availability Zones for 99.99% uptime
- Amazon RDS PostgreSQL: primary + read replicas with automated Multi-AZ failover; FTI in isolated encrypted repository
- ElastiCache (Redis): session management and validation caching; reduces DB load during peak filing periods
- S3 + CloudFront CDN: static asset delivery with geographic caching; offloads 90% of traffic from origin
- AWS WAF + Shield Advanced: DDoS protection, bot mitigation, and OWASP rule enforcement at the edge
- VPC with private subnets: all internal services communicate via private endpoints — zero public internet exposure
- All data in transit: TLS 1.3 + FIPS-validated cryptography; at rest: AWS KMS customer-managed keys

Deployment & Release Strategy

- Infrastructure as Code: Terraform + AWS CloudFormation — every environment (dev, test, staging, PATS, prod) reproducible and version-controlled
- GitLab CI/CD: fully automated build, SAST/DAST scan, test, and deploy pipeline — no manual deployment steps at any stage
- Blue/green deployments: zero-downtime releases with instant rollback by traffic shift to prior environment
- Separate isolated environments: Development, Test, PATS (third-party sandbox), Staging, Production — each independently deployable
- Continuous ATO (cATO): security controls validated in pipeline; no security gate requires manual sign-off to release
- Deployment frequency: multiple times per sprint; production changes coordinated with DOL change management process
- Rollback SLA: any production deployment rolled back within 5 minutes via automated traffic shift if health checks fail
- NACI-cleared staff maintain all production access; least-privilege IAM roles; all access logged and auditable

A5. Scalability & Reliability

Auto-Scaling Architecture

- Auto-scaling triggers at 70% CPU across EKS node groups, RDS read replicas, and SQS consumers — zero manual intervention
- A5. Scalability & Reliability
- CloudFront CDN offloads 90% of static traffic; dramatically reduces origin load during peak filing seasons
- Amazon SQS intake layer absorbs submission spikes without any degradation to the filer-facing API
- RDS read replicas handle concurrent validation queries without write contention during peak periods
- 99.99% uptime SLA: RTO < 1 hour; RPO < 15 minutes; automated Multi-AZ failover with no manual steps
- Multi-AZ deployment: loss of any single Availability Zone causes zero service disruption or data loss
- Chaos engineering practices test failure modes in staging before every peak season deployment

Peak Season Readiness

- Columbus Day peak filing period: load-tested annually with JMeter/Locust simulating maximum concurrent filer volume
- New Year's Day form rollover: tested and executed with zero downtime; contact center staffed for filer questions
- Staffing surge plan: predictive workforce management scales contact center agents ahead of peak filing deadlines
- Performance baselines re-established each sprint; degradation detected before it reaches production
- CloudWatch dashboards: real-time visibility into queue depth, processing latency, error rates, and scaling events
- PagerDuty integration: on-call ops team alerted within 60 seconds of any threshold breach or anomaly
- Quarterly DR drills: full failover exercised and documented; results reviewed with DOL
- Public-facing SLA transparency: DOL can view uptime and performance metrics via shared CloudWatch dashboard

A6. High-Volume Processing (1M+ Filings)

Capacity & Architecture Design

- System sized for 1.4M annual filings with burst headroom — well above the 1M DOL requirement
- SQS message queuing absorbs any submission spike without degrading the intake API — filers never experience slowdowns
- Step Functions process each filing independently and in parallel — not sequentially; throughput scales linearly with load
- Multi-form, multi-year concurrent processing: Form 5500, 5500-SF, 5500-EZ, and all schedules handled simultaneously
- API rate limiting and throttling protect system integrity while maximizing legitimate submission throughput
- Bulk filing API accepts batch submissions — third parties submit thousands of filings per call with structured responses
- Processing workers auto-scale horizontally: each EKS pod handles multiple concurrent filing workflows
- Database read replicas absorb validation query load — no write contention during peak concurrent submission periods

Proven Scale — Satsyil Track Record

- USPTO EDS Platform (Satsyil-built): 20M+ transactions per day — EFAST2's 1.4M annual filings is <1% of our demonstrated daily capacity
- NARA NAC Platform (Satsyil-built): 300M+ digital records processed, validated, and publicly served with 99.8%+ uptime
- Both platforms operate 24x7 with automated deployments, zero-downtime releases, and proactive monitoring
- Annual load testing: JMeter and Locust simulate Columbus Day peak — system must pass before deployment is approved
- Capacity planning review conducted annually ahead of each filing season in coordination with DOL
- Queue depth monitoring: CloudWatch alarms trigger auto-scaling before filers experience any degradation
- Cost-efficient at scale: auto-scaling scales DOWN during off-peak periods — DOL pays only for capacity used
- Throughput SLA: 95th percentile filing acknowledgment issued within 30 seconds of submission receipt at peak load

A7. Quality Controls & Exception Handling

Queuing & Reliability Controls

- Amazon SQS intake queue: every submission durably queued the moment it is received — no filing is lost even if all downstream services are temporarily unavailable
- Exactly-once delivery: SQS + Step Functions use idempotency keys — each filing is processed once and only once; no duplicates, no skipped filings
- Automated retries with exponential backoff: transient failures (network blips, DB timeouts) self-heal automatically without any human intervention
- Dead-letter queues (DLQs): after configurable retry exhaustion, failed filings routed to DLQ; ops team alerted within 60 seconds with filing ID, failure reason, and processing stage
- Pre-submission validation: malformed or non-compliant submissions rejected before entering the processing queue — dramatically reduces exception volume at source
- Duplicate detection: filing fingerprint hashing prevents double-processing of the same submission; idempotent resubmission supported
- End-to-end traceability: every filing has a unique tracking ID from intake through distribution; every status transition logged with timestamp and audit record

Exception Handling Workflow

- Exception triage: DLQ alerts include structured metadata — filing ID, submitter, form type, failure code, processing stage — enabling rapid root cause analysis
- Automated remediation: common exception classes (schema version mismatch, temporary downstream outage) trigger automated resolution workflows without ops team involvement
- Manual remediation: ops team reviews DLQ entries; corrective action documented in Jira with DOL visibility; filing reprocessed once issue resolved
- Filer notification: if a filing cannot be processed after all retries, the filer receives a structured error notification with specific guidance for resubmission
- Quality control gates: automated validation checks at each pipeline stage prevent downstream propagation of bad data
- Operational metrics reviewed weekly: exception rates, DLQ volume, retry frequency, and resolution time tracked against SLA targets
- Continuous improvement: recurring exception patterns trigger root cause analysis and permanent system fix — not just one-time remediation
- Zero silent failures: every filing that enters the system either reaches a successful distribution state or generates a tracked, alerted,

B. Security and Vulnerability Management

B1. Authentication & Authorization

Filer & Government Authentication

- Login.gov via OIDC/SAML 2.0 for all public filers at IAL-1/IAL-2 and AAL2 — federally mandated identity standard
- PIV/HSPD-12 smart card authentication for all government users via AWS IAM Identity Center federation
- MFA enforced on every session — no single-factor access to any system component
- Short-lived JWT tokens with automatic expiry; continuous session monitoring flags anomalous behavior
- Role-based access control (RBAC): DOL staff, contractor staff, filers, and third parties each have scoped, least-privilege roles
- Admin panel access requires DOL-approved credentials with full audit trail on every configuration change
- Government user accounts provisioned through DOL identity management with mandatory periodic recertification
- Zero standing access: all privileged access time-limited and requires just-in-time approval workflow

Third-Party & API Authorization

- Third-party software providers issued scoped API keys after successful PATS certification only — not before
- API keys carry minimum necessary permissions; cannot be elevated without new certification cycle
- NIST SP 800-63 credential lifecycle management: rotation, revocation, and audit logging for all API credentials
- API key revocation takes effect within seconds — immediate response to any compromise or compliance violation
- All authentication events logged to tamper-evident audit store; available for DOL review at any time
- Webhook authentication: all outbound callbacks signed with HMAC-SHA256 — recipients can verify Satsyil origin
- OAuth 2.0 client credentials flow for machine-to-machine integrations (IRS, PBGC); no shared secrets in transit
- Quarterly access reviews: all active API keys and user accounts reviewed and recertified with DOL participation

B2. Protection of Sensitive Data

Encryption & Data Minimization

- All data in transit encrypted with TLS 1.3 using FIPS 140-2 validated cryptographic modules — no exceptions
- All data at rest encrypted with AWS KMS customer-managed keys (CMKs); keys rotated annually
- Form 5558 FTI data stored in dedicated, isolated repository — access controlled, encrypted, and audited separately per IRS Pub. 1075
- Data minimization: only fields required by the DER are collected; no additional PII captured beyond regulatory mandate
- PII/FTI automatically detected and redacted from all logs, error messages, and API responses before storage
- Supervised ML models (SageMaker) identify and mask PII in submission data prior to public release
- Database encryption at column level for all sensitive fields (SSN, EIN, account numbers)
- S3 bucket policies enforce server-side encryption (SSE-KMS) and block all public access

FTI & Regulatory Compliance

- IRS Publication 1075 compliance: FTI data in fully segregated, audited environment with access limited to authorized personnel
- NIST SP 800-53 Rev. 5 controls implemented and continuously validated via automated compliance scanning
- FedRAMP Moderate authorization maintained; all new services evaluated against FedRAMP control baseline before deployment
- Privacy Impact Assessment (PIA) maintained and updated with any new data element or processing change
- Data retention policies enforced automatically: records purged per DOL/NARA retention schedule
- Cross-region replication disabled for FTI data — data sovereignty enforced within approved AWS GovCloud regions only
- Secure deletion: data purged using NIST 800-88 compliant methods; deletion confirmed with audit record
- Annual third-party penetration test targeting data protection controls; findings remediated before next filing season

B3. Audit Logging & Traceability

Comprehensive Audit Coverage

- Every submission logged from intake through distribution: filing ID, submitter, timestamp, processing stage, and outcome
- Every authentication event logged: login, logout, failed attempts, MFA completions, token issuance and expiry
- Every admin configuration change logged: who made it, what changed, when, and government approval reference
- Every API call logged with caller identity, endpoint, payload hash, response code, and latency
- Every file access to FTI repository logged with user ID, timestamp, file identifier, and access purpose
- Every downstream delivery logged: file sent to IRS, PBGC, DOL, NARA with delivery confirmation and receipt timestamp
- Dead-letter queue entries logged with full failure context — no silent failures, every exception is traceable
- All logs are tamper-evident: CloudTrail + CloudWatch Logs with integrity validation; log deletion triggers immediate alert

DOL Access & Retention

- DOL has real-time read access to all audit logs via shared CloudWatch dashboards — no access request required
- Logs retained for minimum 3 years in hot storage; 7 years in cold storage per federal records requirements
- Log search and filtering: DOL staff can query by filing ID, filer EIN, date range, or event type within seconds
- Automated anomaly detection: ML models flag unusual access patterns, bulk downloads, or off-hours admin activity
- Audit report generation: automated monthly audit summaries delivered to DOL in structured format
- SIEM integration: logs forwarded to DOL SIEM if required; standard format (CEF/JSON) supported
- Incident response: all security events traceable end-to-end; forensic log packages generated within 4 hours of request
- Log integrity: AWS CloudTrail log file validation enabled; any tampering detected and alerted automatically

B4. Secure Document & Attachment Handling

Secure Intake & Scanning

- Every attachment scanned for malware and viruses using AWS-native scanning plus third-party AV engine on upload
- File type validation: only permitted formats (PDF, XML, CSV per DER) accepted; all others rejected with structured error
- File size limits enforced at API gateway and upload layer; oversized submissions rejected before scanning
- Attachments stored in isolated S3 bucket separate from filing metadata — never co-mingled with public data
- All attachments encrypted at rest with KMS CMK; unique key per document class (FTI vs. non-FTI)
- Attachment access requires authenticated submission context — documents cannot be accessed without valid filing reference
- Attachment integrity: SHA-256 hash computed on upload and verified on every retrieval — any tampering detected
- FTI-containing attachments automatically flagged and routed to IRS Pub. 1075-compliant isolated handling pipeline

Distribution & Lifecycle

- Attachments transmitted to downstream agencies (IRS, PBGC, DOL) via encrypted SFTP with file-level hash verification
- NARA archive: attachments packaged in non-proprietary formats with metadata; integrity verified before delivery
- Public dissemination: PII/FTI stripped from attachments before any public release; ML-verified and human-reviewed
- Document retention: attachments retained per DOL/NARA schedule; automated deletion with audit record at end of retention period
- Resubmission handling: updated attachments replace prior versions in controlled manner; prior versions archived with audit trail
- Failed attachment delivery: DLQ captures delivery failures; ops team alerted with document ID and failure context
- Attachment audit: every access, download, or transmission logged with user/system identity and timestamp
- Secure deletion: when retention period expires, attachments deleted using NIST 800-88 guidelines; confirmed in audit log

B5. Vulnerability Monitoring, Patching & Remediation

Continuous Monitoring & Detection

- AWS GuardDuty: continuous threat detection; ML-based anomaly detection identifies unusual API calls, network flows, and access patterns
- AWS Security Hub: aggregates findings from GuardDuty, Inspector, Config, and Macie into single prioritized dashboard
- AWS Config: continuous configuration compliance monitoring; drift from approved baseline triggers immediate alert
- SAST scanning on every code commit in GitLab CI/CD — no code reaches staging without passing security gates
- DAST scanning against staging environment before every production deployment
- Automated dependency scanning (Dependabot/Snyk): known CVEs in third-party libraries flagged within hours of disclosure
- Container image scanning (Amazon ECR): every image scanned before deployment; critical CVE images blocked from EKS
- Annual third-party penetration test plus quarterly vulnerability assessments; results shared with DOL within 5 business days

Patching SLAs & Remediation

- Critical CVE (CVSS 9.0+): patched within 15 days of disclosure — no exceptions, tracked to closure in Jira with DOL visibility
- High CVE (CVSS 7.0-8.9): patched within 30 days; compensating controls documented if patch requires coordination
- Medium CVE (CVSS 4.0-6.9): patched within 90 days as part of regular sprint cycle
- OS and runtime patching: automated via AWS Systems Manager Patch Manager; compliance reports available to DOL
- Zero-day response: incident response activated within 1 hour of credible zero-day disclosure; DOL notified same day
- Continuous ATO (cATO): security controls validated in CI/CD pipeline; ATO maintained continuously without annual re-accreditation
- POA&M management: all open findings tracked with target dates; reviewed monthly with DOL
- Security metrics reported monthly to DOL: open findings by severity, mean time to patch, scan coverage, and exception count

C. Data Quality and Validation

C1. Field-Level & Cross-Field Validations

Field-Level Validation Engine

- Schema validation against machine-readable DER (JSON/XML) artifacts on every submission before data enters processing queue
- Field-level checks: data type, format, length, allowable values, and required-field presence validated for every element
- Numeric range validations: contribution amounts, participant counts, and asset values checked against regulatory thresholds
- Date validations: plan year dates, filing deadlines, and fiscal year boundaries enforced with form-year-aware logic
- EIN and plan number validation: format checks plus cross-reference against known entity registry where applicable
- Schedule attachment validation: required schedules enforced based on plan type, participant count, and funding status
- Configurable validation engine: DOL staff add or modify field-level rules in admin panel without code deployment
- Same validation engine used in PATS sandbox and production — third parties test against identical rules before going live

Cross-Field Validation Rules

- Cross-field consistency: total assets must reconcile across schedules; contributions cannot exceed plan limits; participant counts cross-validated
- Form-type-specific rule sets: Form 5500, 5500-SF, and 5500-EZ each carry distinct cross-field rules enforced independently
- Plan characteristic validations: large vs. small plan rules, funding method requirements, and exemption eligibility cross-checked
- Schedule interdependency: Schedule H totals must align with Schedule A entries; discrepancies flagged before acceptance
- Historical consistency: current-year values cross-referenced against prior-year filing to flag implausible changes
- Entity validation: plan sponsor EIN cross-validated against submitted employer identification data; mismatches flagged
- All cross-field failures return structured error codes with specific field references — filers know exactly what to fix
- Pre-submission validator available via UI and API: filers and third parties run all validation checks before formal submission

C2. Business Rule Enforcement

Rule Engine Architecture

- Business rules stored as versioned metadata (JSON rule definitions) — not hardcoded; rules modified in admin panel without code deployment
- Edit-test logic engine executes rules against submission data in real time; rules fire in defined sequence with dependency management
- Form-year tagging: each submission matched to its applicable rule set by form year — no risk of applying wrong-year rules
- Rule versioning: every rule change versioned with effective date, author, and government approval reference
- Configurable severity levels: rules classified as fatal (rejection), warning (advisory), or informational — DOL controls classification
- Rule conflict detection: admin panel flags conflicts between rules before they are published to production
- PATS test harness updated with new rules before each filing season; third parties validate against updated rules in sandbox
- Rule audit trail: every rule execution logged with submission ID, rule ID, outcome, and timestamp — fully auditable by DOL

Regulatory Alignment

- Rules derived directly from DOL DER artifacts; traceability matrix maps each rule to its regulatory source
- Annual rule review: Satsyil and DOL jointly review rule set before each filing season; changes approved before rollout
- IRS coordination: Form 5558 rules aligned with IRS requirements; rule changes triggered by IRS publications incorporated within 30 days
- PBGC coordination: premium filing rules aligned with PBGC annual guidance; Satsyil participates in annual inter-agency rule review
- Regulatory change management: when DOL issues updated guidance, rule changes implemented and tested before effective date
- Exception reporting: rules suspended or waived by DOL documented with authority and effective period; automatically reinstated
- Rule performance monitoring: rules generating high false-positive rates flagged for DOL review; accuracy metrics tracked monthly
- Documentation: every active rule documented with

C3. Error Messaging & Correction Workflow

Structured Error Responses

- Every validation failure returns a structured error object: error code, field path, rule ID, plain-language description, and correction guidance
- Error severity classification: fatal errors block submission; warnings allow submission with notation; informational messages advisory only
- Error messages written in plain language per the Plain Language Act — no cryptic codes without explanation
- Field-level error highlighting in IFILE portal: errors shown inline next to the specific field, not just in a summary list
- Bulk filing API error responses include structured JSON with per-record error arrays — third parties process errors programmatically
- Error grouping: related errors grouped by schedule/section so filers address them systematically
- Error persistence: errors saved to filer dashboard so they can return and correct without re-entering the entire form
- Error metrics: DOL receives monthly report on most common error codes, error rates by form type, and correction success rates

Correction & Resubmission Workflow

- Filers can correct and resubmit without restarting the entire filing workflow — prior valid entries preserved
- Amendment workflow: superseded filings tracked and linked to amended filings; DOL can view full amendment history for any plan
- Pre-submission correction: validation available as a dry-run before formal submission — filers fix all errors before committing
- Partial save: IFILE portal allows saving progress and returning; incomplete filings auto-expire per DOL policy with notification
- Third-party correction flow: API returns errors with filing reference ID; third party corrects and resubmits using same reference
- Contact center integration: error notifications include option to contact support; agent sees same error details as filer
- Duplicate detection: resubmitted corrected filings matched to original via plan/EIN/year key — no duplicate processing
- Correction audit trail: every correction and resubmission logged with timestamp, changed fields, and submitter identity

C4. Pre-Submission Validation

Pre-Submission Tools

- Pre-submission validator available via IFILE web portal and REST API — filers and third parties run complete validation before formally submitting
- Same validation engine as production: pre-submission validator uses identical schema, rules, and business logic — no surprises at submission
- Dry-run mode: complete validation executed without creating a formal submission record — safe to test repeatedly
- Pre-population of prior-year data: IFILE pre-fills fields from the prior accepted filing; filer reviews and updates changed values only
- Third-party pre-validation API: software providers integrate pre-validation into their workflow so filers see errors before export to EFAST2
- PATS sandbox provides production-identical pre-validation for third-party certification testing
- Proven impact: analogous pre-validation gates at NARA reduced ingestion failures by 33%
- Pre-validation results cached for 24 hours: filer can review errors, correct, and resubmit without re-running unchanged

Quality Gates Before Acceptance

- Malformed submissions rejected at intake layer before entering processing queue — SQS never receives invalid submissions
- Schema validation is the first gate: XML/JSON must be well-formed and schema-compliant before any business rule fires
- Business rule gate: all fatal rules must pass before submission is accepted; advisory rules generate warnings, not rejection
- Entity validation gate: EIN and plan number must resolve to a recognized entity before acceptance
- Attachment gate: all required attachments present, file type valid, virus scan clean, and hash verified before acceptance
- Certification gate for third parties: software provider API key must be active and certified for the form type being submitted
- Duplicate gate: fingerprint hash checked against recent submissions; exact duplicates rejected with reference to original
- Every rejected submission receives a structured rejection notice with all failing checks listed — no partial information

C5. Automated Acknowledgments & Status Responses

Acknowledgment Architecture

- Acknowledgment issued within seconds of submission receipt — before any processing begins; includes unique tracking ID
- Acknowledgment delivered via API response (synchronous), email notification, and filer dashboard update simultaneously
- Acknowledgment content: filing reference ID, submitter identity, form type, plan year, timestamp, and initial acceptance status
- Processing status updates delivered at each pipeline stage: received, validating, processing, distributing, complete
- Filer dashboard: real-time status for all submissions; sortable by status, date, form type, and plan number
- Third-party API polling: status endpoint returns current processing state for any filing by reference ID
- Webhook push: third parties register for push notifications; status updates delivered within 30 seconds of each transition
- Error acknowledgments: if validation fails post-acceptance, structured error notification sent with all failing checks and correction guidance

Status Codes & Error Responses

- Standardized status code taxonomy: RECEIVED, VALIDATING, ACCEPTED, PROCESSING, DISTRIBUTED, REJECTED, ERROR — consistent across API and UI
- Rejection notices include: rejection code, all failing validation checks with field references, suggested corrections, and contact information
- Distribution confirmation: once filed data reaches IRS, PBGC, and DOL, confirmation status updated with delivery timestamp
- NARA archive confirmation: annual archive delivery confirmed with package ID and delivery receipt in filer record
- Amended filing acknowledgments: amendments acknowledged separately; linked to original filing with supersession notation
- Contact center linkage: acknowledgment emails include support contact; agents see same status in CRM when filer calls
- Third-party error response format: JSON error array with per-record status for bulk submissions — each record acknowledged individually
- All acknowledgments and status events retained for 7 years

D. Annual Rollover and Version Changes

D1. Adding or Modifying Fields & Validations

How Fields Are Added or Modified

- Fields defined in machine-readable DER (JSON/XML) metadata — not hardcoded; new fields added by updating the DER artifact in admin panel
- Admin configuration panel: authorized DOL staff add a new field by defining name, type, format, rules, and effective date — no developer involvement
- Field addition triggers automatic: IFILE UI update, API schema update, PATS test case generation, and documentation update
- Effective date control: new fields scheduled to activate on a specific date (e.g., January 1) with no deployment required on that date
- Backward compatibility: new required fields carry a grace period for existing third-party integrators; optional in first season, required in subsequent
- Validation rule modification: changing a field's rules follows same admin panel workflow with government approval gate before publishing
- Testing: every field addition or change automatically generates regression test cases; CI/CD pipeline runs them before any deployment
- Audit trail: every field change recorded with who added it, what the change was, government approval reference, and effective date

Impact Analysis & Notification

- Impact analysis: admin panel shows how many prior-year submissions would be affected by a new rule — DOL calibrates severity before activating
- Rule testing: all modified rules run against historical submission corpus to measure impact before activation
- Third-party notification: 90-day advance notice of all field or rule changes sent to all registered developer portal subscribers
- PATS update: test harness updated with new fields and rules before notification goes out — developers can test immediately
- Developer conference: annual conference held to walk through all upcoming field and rule changes before new filing season begins
- Documentation updated: every field change reflected in API docs, schema files, and developer portal simultaneously with activation
- Rule deprecation: old rules retired with notice period; deprecated rules flagged in documentation with effective-end date
- Rollback: field additions can be deactivated via admin panel if issues discovered; rollback documented and DOL notified

D2. Annual Rollover & Version Changes

Annual Rollover Execution

- New Year's Day rollover: form year activation scheduled and tested in staging months in advance; executes automatically at midnight with zero downtime
- Blue/green deployment: new form year activates by traffic shift — instant rollback to prior year if any issue detected in first 30 minutes
- Form year coexistence: prior-year forms remain active for amended filings; new-year forms activate for current-year submissions simultaneously
- PATS updated: new form year available in sandbox 90 days before activation so third parties can test and certify against new schemas
- Contact center staffed: agents briefed on all changes before January 1; knowledge base updated with new form year FAQs
- IRS and PBGC coordination: distribution files updated for new form year formats; downstream recipients notified 90+ days in advance
- Third-party developer notification: change summary distributed 90 days before rollover; API change log published in developer portal
- Columbus Day peak filing: annual load test conducted before October to validate system handles peak volume with new form year schemas

Version Management

- Every form version (5500, 5500-SF, 5500-EZ) independently versioned; multiple form years coexist in production simultaneously
- Semantic versioning for APIs: major version = breaking change; minor = additive; patch = fix; clients declare which version they target
- API version support window: minimum 2 prior versions supported concurrently; deprecation announced 12 months in advance
- Schema versioning: XML and JSON schemas versioned with namespace and version identifier; validation engine routes submission to correct version automatically
- Documentation versioned: developer portal maintains changelog and migration guides for every version transition
- Rollback capability: any version activation can be reversed within 15 minutes by shifting traffic back to prior environment
- Version audit: every submission tagged with form version, schema version, and API version — queryable by DOL for compliance analysis
- Version sunset: 18-month sunset period with DOL approval when a form version reaches end-of-life; all filers migrated before decommission

D3. Configuration-Driven & Metadata-Driven Updates

Configuration Architecture

- All form definitions, field metadata, validation rules, and business logic stored as versioned configuration — not in application code
- Configuration stored in version-controlled GitLab repository; every change peer-reviewed and DOL-approved before merge
- Metadata-driven rendering: IFILE portal renders form fields dynamically from DER metadata — no UI code changes required to add or modify fields
- Rule engine reads configuration at runtime: no restart or redeployment needed to activate a rule change scheduled for a future date
- Admin configuration panel: authorized DOL staff manage form definitions, rules, and schedules directly without requiring developer involvement
- Environment promotion: configuration changes flow dev → test → staging → production with automated validation at each gate
- Feature flags: new configuration features activated per-environment and per-user-segment before full rollout — safe, gradual activation
- Configuration audit: every configuration state is reproducible; any prior configuration can be restored from version history within minutes

Benefits & Proof Points

- Zero code deployment for routine annual updates: field additions, rule changes, and schedule modifications require no developer sprint
- DOL retains control: government staff preview, approve, and schedule configuration changes without waiting for contractor release cycle
- Rapid response: when DOL issues updated guidance, configuration updated and activated within days — not weeks
- Proven at scale: Satsyil's configuration-driven approach at NARA supported hundreds of metadata schema updates without application redeployment
- Cost efficiency: eliminating developer effort for routine annual changes reduces O&M cost; capacity focused on enhancements
- Risk reduction: configuration changes isolated from application code — a misconfigured rule cannot crash the application
- Transparency: all active configuration visible to DOL in admin panel; no hidden logic, no undocumented behavior
- Disaster recovery: if configuration corruption occurs, prior version restored from GitLab history within 15 minutes

D4. Minimizing Code Changes & UI Redeployment

How We Minimize Code Changes

- Form UI rendered dynamically from DER metadata: adding a new field to the IFILE portal requires no UI code change — only a metadata update
- Validation engine is data-driven: all validation logic defined in configuration; new rules added without touching application code
- API schema auto-generated from DER artifacts: when metadata changes, API schema updates automatically — no manual API code changes
- Documentation auto-generated from DER: developer portal schema documentation regenerated on every metadata update — always in sync
- PATS test cases auto-generated from rule changes: new rules automatically produce test scenarios in the PATS harness
- Scheduled activations: configuration changes staged weeks in advance and activate automatically on target date — no deployment-night intervention
- GitLab CI/CD: when code changes are required, fully automated pipeline completes deployment in under 30 minutes with zero downtime
- Blue/green deployments: even when code changes are needed, zero downtime guaranteed; traffic shifted after health check passes

Deployment Guardrails & Metrics

- Target: zero code deployments required for annual form rollover — all routine changes handled through configuration pipeline
- Automated regression: every configuration change triggers full regression test suite before activation — no manual testing required
- Deployment frequency: code deployments occur on sprint cadence (bi-weekly); configuration updates can deploy daily if needed
- Rollback SLA: any deployment (code or configuration) rolled back within 5 minutes if health checks fail post-activation
- Change advisory board: all production deployments reviewed in weekly CAB meeting with DOL participation before scheduling
- Deployment blackout windows: no deployments during peak filing periods (Oct 1-15, Jan 1-7) without DOL emergency approval
- Deployment audit: every deployment logged with change description, approver, deployment time, and post-deployment validation results
- Continuous improvement: deployment metrics (frequency, duration, rollback rate) reported to DOL monthly; targets reviewed annually

E. AI Use in Development, PII Identification, and Testing

E1. AI in Code Development

AI-Assisted Development Tools

- GitHub Copilot integrated into developer IDE: AI-assisted code completion, boilerplate generation, and in-context documentation drafting
- AI code review: automated PR review tools flag code smells, security anti-patterns, and style violations before human review
- LLM-assisted API documentation: OpenAPI specs and developer portal docs drafted by AI from code annotations; reviewed before publishing
- AI-assisted configuration generation: new DER-based rule configurations drafted by AI from regulatory text; DOL-approved before activation
- AI-powered semantic code search helps developers locate relevant implementation patterns across the large codebase quickly
- Automated boilerplate: microservice scaffolding, test class stubs, and migration scripts generated by AI — reduces developer setup time by 60%+
- AI-assisted code refactoring: legacy code analysis tools identify refactoring opportunities; AI generates candidate refactors for human review
- All AI-generated code reviewed by a human engineer before merge — no AI code enters production without engineer approval

Governance & Controls

- No federal submission data used for AI model training — all AI tools operate on synthetic or anonymized datasets only
- Approved tool list: only vetted AI development tools permitted; unapproved tools blocked by developer workstation policy
- Code provenance: AI-assisted code flagged in commit history; auditors can identify AI-assisted vs. human-written code
- Security review: AI-generated code undergoes same SAST/DAST scanning as human-written code — no exemptions
- No AI generates security-sensitive code (cryptography, authentication, key management) without mandatory expert human review
- Vendor agreements: AI tool vendors contractually prohibited from training on Satsyil or DOL code or data
- Quarterly review: AI tool usage, code quality metrics, and defect rates reviewed quarterly; tools removed if metrics degrade
- Transparency: DOL informed of all AI development tools in use; tool list maintained in project management documentation

E2. AI in Testing & Test Automation

AI-Powered Test Generation

- LLM-driven test case generation: AI analyzes DER requirements and business rule definitions to auto-generate comprehensive test scenarios
- AI-generated edge case discovery: AI models identify boundary conditions and unusual input combinations that human testers commonly miss
- PATS test harness: AI generates certification test scenarios from rule changes automatically — new rules ship with test coverage on day one
- Load test script generation: AI generates JMeter and Locust scripts from API specification; humans review and tune parameters
- Regression suite expansion: AI proposes new regression tests when production defects are fixed — prevents defect recurrence
- Test data synthesis: AI generates realistic synthetic Form 5500 data covering all form types, plan sizes, and edge cases — no real PII used
- Visual regression testing: AI compares IFILE portal screenshots across deployments to detect unintended UI changes
- Section 508/WCAG 2.2 AA: automated accessibility testing on every deployment using axe-core and DHS Trusted Tester methodology

Coverage & Quality Metrics

- Target: 85%+ code coverage across all microservices; AI-generated tests have increased coverage by 40%+ vs. manual testing alone
- Flaky test detection: AI identifies tests with inconsistent results; flaky tests quarantined and fixed before they mask real defects
- Test gap analysis: AI reports on areas of code not covered by tests; gaps reviewed and prioritized in sprint planning
- Defect prediction: ML model trained on historical defect data predicts which code changes are highest risk; flagged for additional review
- Quality gate: CI/CD pipeline blocks deployment if coverage drops below threshold or AI risk score exceeds limit
- Performance baseline: AI monitors test execution times and flags tests that degrade significantly
- Mutation testing augmented by AI: AI identifies which mutations are most likely to reveal logic gaps; focuses human review on high-risk areas
- Test results reported to DOL monthly: coverage percentage, defect rates, AI-generated vs. manual test ratio, and trend analysis

E3. AI for Quality Improvement

Operational Quality AI

- Amazon SageMaker anomaly detection: ML models monitor submission patterns in real time; unusual volumes, error spikes, or processing delays flagged instantly
- Predictive auto-scaling: ML models forecast filing volume based on historical patterns and calendar events; EKS scales ahead of demand
- AI-driven error pattern analysis: recurring error codes analyzed to identify systemic filer issues; DOL informed for potential outreach or guidance updates
- Rejected submission analysis: AI identifies patterns in rejected filings to recommend validation rule improvements that reduce filer burden
- Contact center quality: AI analyzes call transcripts to identify common filer pain points; findings inform IFILE UX improvements
- Data quality scoring: ML model scores each submission on predicted data quality risk; high-risk submissions flagged for enhanced review
- PII detection: supervised ML models detect PII/FTI in submissions, logs, and API responses before public release — proven at USPTO
- Performance degradation detection: AI monitors API latency distributions and flags degradation before it reaches user-visible thresholds

Improvement Feedback Loops

- Monthly AI quality report to DOL: anomalies detected, error patterns identified, scaling events, and PII detection accuracy metrics
- Closed-loop improvement: AI findings feed back into validation rule updates, IFILE UX improvements, and contact center training
- Model drift monitoring: AI model performance tracked over time; models retrained when accuracy drops below threshold
- Human review of AI recommendations: all AI quality recommendations reviewed by Satsyil quality team before any system change
- Filer experience improvement: AI tracks form completion rates, abandonment points, and error frequency to guide UX enhancements
- Rule effectiveness measurement: AI measures which validation rules generate the most corrections vs. disputes; informs rule refinement
- Third-party quality tracking: AI monitors error rates by software provider; providers with consistently high error rates flagged for outreach
- Continuous model improvement: models updated quarterly using new data; prior model versions archived for comparison and rollback

E4. AI for Documentation & Defect Detection

AI-Assisted Documentation

- API documentation auto-generated from OpenAPI annotations; developer portal always synchronized with deployed API version
- DER-to-plain-language: AI drafts plain-language field descriptions from technical DER specifications; reviewed by Satsyil and DOL before publishing
- Release notes generation: AI summarizes code changes into human-readable release notes; reviewed by project manager before distribution
- Knowledge base articles: AI drafts contact center knowledge base articles from ticket resolutions; reviewed by subject matter expert before publishing
- Test documentation: AI generates test case documentation from test code; keeps test library documentation current without manual effort
- Runbook generation: AI drafts operational runbooks from system architecture and incident history; reviewed and approved by ops lead
- Regulatory change tracking: AI monitors Federal Register and DOL publications for EFAST2-relevant changes; flags potential impacts for human review
- All AI-generated documentation reviewed and approved by a human before publication — AI drafts, humans certify

AI Defect Detection

- Static analysis: AI-powered SAST tools detect security vulnerabilities, logic errors, and code quality issues on every commit
- Log anomaly detection: AI analyzes application logs in real time to detect error patterns that indicate latent defects before user impact
- Production defect prediction: ML model trained on historical defects predicts which code areas are highest-risk before each release
- Database query analysis: AI monitors slow query logs and flags queries that could cause performance issues under peak load
- Dependency vulnerability monitoring: AI tracks CVE feeds and alerts within hours of a vulnerability affecting any dependency in the stack
- Configuration drift detection: AI compares live infrastructure state against approved baseline; drift detected and alerted within minutes
- API contract testing: AI monitors API response structures against documented contracts; schema violations detected before reaching filers
- All AI defect findings triaged by Satsyil engineering team; findings rated by severity and assigned to sprint with DOL visibility in Jira

E5. AI Governance & PII Identification Controls

AI Governance Framework

- Satsyil AI Governance Framework: documented policy governing all AI tool selection, approval, deployment, monitoring, and decommission
- AI inventory: all AI tools and models in use documented with purpose, data inputs, risk classification, and approval status
- No AI adjudicates filing acceptance: all consequential decisions (rejection, hold, referral) require human review — AI provides recommendations only
- Human-in-the-loop mandatory for: filing rejection, PII redaction decisions, security incident classification, and regulatory interpretation
- Model governance: all ML models versioned, documented, and validated before production deployment; prior versions retained for rollback
- Bias monitoring: AI models monitored for disparate impact across filer types; bias findings reviewed quarterly
- AI risk classification: each AI use case assigned risk level (low/medium/high); high-risk uses require DOL awareness and additional controls
- Federal AI policy compliance: all AI use aligned with OMB AI

PII Identification & Control

- Supervised ML models trained on federal filing data detect PII (SSN, name, address, account numbers) in submissions, logs, and API responses
- Proven at USPTO: Satsyil's PII detection models achieve 98%+ recall on federal filing data; false negatives reviewed manually
- Automatic redaction: detected PII in non-required fields redacted before data enters public dissemination pipeline
- FTI-specific controls: Form 5558 FTI data handled under separate ML pipeline with additional controls per IRS Pub. 1075
- No submission data used for model training: PII models trained only on DOL-approved synthetic datasets — no real filer data ever used
- PII detection audit: every detection event logged with field, detected entity type, action taken, and confidence score
- False positive management: PII model false positives reviewed by human analyst; corrections fed back to improve model accuracy
- Annual PII control review: DOL reviews PII detection

F. Data Migration and Transition Planning

F1. Large-Scale Data & Document Migration

Migration Architecture

- AWS Glue / Databricks ETL pipelines: fully automated, parallelized migration with complete audit logging at every stage
- Legacy EFAST2 data AND document/attachment migration: all filed Form 5500 attachments migrated with SHA-256 hash integrity checks
- Migration phased by data vintage: most recent years migrated first (highest query demand); historical data follows in subsequent phases
- No service interruption: legacy EFAST2 continues serving filers during migration; new system activated only after migration validation completes
- Structured data migration: filing records, plan data, filer profiles, and status histories migrated with full relational integrity
- Unstructured document migration: PDF, XML, and binary attachments migrated to S3 with metadata tagging and hash verification
- FTI data migration: Form 5558 FTI data migrated under separate, audited pipeline meeting IRS Pub. 1075 requirements
- Proven at scale: Satsyil migrated 300M+ digital records for NARA NAC — reducing ingestion failures by 33% and improving data quality by 60%

Migration Execution & Tracking

- Migration dashboard: real-time visibility into records migrated, records remaining, error count, and throughput — accessible by DOL at all times
- Per-record audit trail: every migrated record logged with source ID, target ID, migration timestamp, and validation outcome
- Incremental migration: delta loads capture new submissions made to legacy system during migration period; no data gap at cutover
- Data profiling: automated profiling identifies data anomalies, format inconsistencies, and encoding issues in legacy data before migration
- Legacy format transformation: legacy XML converted to current schema format with transformation rules documented and DOL-approved
- Migration completion criteria: DOL and Satsyil jointly define acceptance criteria (record counts, hash verification rates, spot-check pass rates)
- Parallelized EKS workers scale to migrate millions of records per day; timeline planned with DOL and tracked weekly
- Post-migration reporting: final migration report delivered to DOL with complete statistics, anomalies encountered, and resolution actions

F2. Validation & Reconciliation

Automated Validation Pipeline

- Source-to-target reconciliation: record counts, field values, and checksum totals compared between legacy and new system at every migration batch
- Field-level validation: every migrated field validated against current DER schema; format conversions verified for correctness
- Business rule re-execution: migrated submissions re-run through current validation engine; discrepancies between legacy and current rules documented
- Referential integrity checks: all foreign key relationships (plan to filer to submission to attachment) verified post-migration
- Hash verification: SHA-256 hash of every migrated document compared against source hash; any mismatch flagged and reprocessed
- Statistical sampling: automated spot-check of 5% random sample with 100% manual review of any anomalies found
- Aggregate reconciliation: total asset values, participant counts, and contribution totals reconciled at plan and year level against legacy reports
- FTI reconciliation: Form 5558 record counts and field values

Reconciliation Reporting

- Daily reconciliation reports: automated reports comparing legacy vs. new system record counts by form type and plan year — delivered to DOL each morning
- Exception reports: all records failing reconciliation checks listed with reason code; Satsyil engineering resolves within 48 hours
- Reconciliation dashboard: DOL-accessible live dashboard showing migration completeness percentage by data category and form year
- Independent verification: DOL retains right to perform independent reconciliation queries against both systems during migration period
- Discrepancy resolution log: every reconciliation discrepancy documented with root cause, resolution action, and timestamp
- Pre-cutover sign-off: DOL formally approves reconciliation results before any production traffic shifts to new system
- Post-cutover reconciliation: 30-day post-cutover monitoring period with daily reconciliation reports to confirm ongoing data integrity
- Reconciliation artifacts retained: all reconciliation reports and

F3. Rollback Planning

Rollback Architecture

- Legacy system maintained in read-only state for 90 days post-cutover: filers can query historical data; no new filings accepted by legacy
- Traffic routing via DNS/ALB: cutover implemented by shifting traffic to new system; reversal takes under 5 minutes — DNS TTL pre-shortened before cutover
- Go/no-go gates: formal decision points before each migration phase and before final cutover; DOL participates in all go/no-go reviews
- Pre-defined rollback trigger criteria: >1% reconciliation failure rate, >5% processing error rate, or critical data integrity issue triggers rollback
- Rollback execution: if trigger criteria met, traffic shifted back to legacy within 5 minutes; DOL notified immediately with incident report
- Data rollback: any data written to new system during failed cutover window identified by timestamp and purged; legacy data remains authoritative
- Rollback drill: full rollback procedure rehearsed in staging 30 days before production cutover; DOL observes the drill
- Post-rollback analysis: root cause analysis completed within 24 hours; remediation plan presented to DOL before rescheduling cutover

Phased Cutover Strategy

- Phase 1 — Read traffic: new system serves read queries (status, history) while legacy handles writes; validates new system under real load
- Phase 2 — Non-peak write traffic: new system accepts submissions during off-peak hours; legacy remains primary during peak periods
- Phase 3 — Full cutover: all traffic shifted to new system after Phase 2 validation; legacy in read-only mode
- Each phase has defined success criteria and minimum dwell time before advancing; DOL approves each phase transition
- Parallel operation logging: during phased cutover, submissions processed by both systems in shadow mode to validate consistency
- Filer communication: clear communication plan for each cutover phase; filers informed of behavior changes with advance notice
- Third-party notification: all registered third-party software providers notified of cutover schedule 90+ days in advance with API migration guide
- Emergency contacts: 24/7 escalation path to DOL COTR and Satsyil program manager during all cutover windows

F4. Handling Undocumented Logic & Business Rules

Discovery Approach

- 30-day structured knowledge transfer with incumbent contractor: workshops, code walkthroughs, and documentation sessions before decommission
- AI-assisted code analysis: ML tools parse legacy codebase to surface undocumented conditional logic, edge-case handlers, and embedded business rules
- Historical data mining: AI analyzes 10+ years of submission outcomes to reverse-engineer business rules that exist only in legacy code behavior
- Incumbent interview program: systematic interviews with current operations and development staff to capture tribal knowledge
- Legacy test execution: existing legacy test suites run against new system to reveal behavioral differences requiring investigation
- Shadow mode validation: during parallel operation, new system results compared against legacy results for identical inputs; differences investigated
- Filer impact analysis: edge-case rules affecting small filer populations identified and verified with DOL subject matter experts
- DOL SME engagement: DOL EBSA program staff consulted on undocumented rules discovered during analysis — authoritative source on regulatory intent

Documentation & Verification

- Business rule registry: all discovered rules documented with source (code, incumbent interview, historical analysis), regulatory basis, and test cases
- Gap analysis report: delivered to DOL within 60 days of contract award; lists all undocumented rules discovered and proposed treatment
- DOL approval required: every undocumented rule either confirmed by DOL as valid or identified as a legacy defect to be corrected
- Test case coverage: every discovered rule covered by automated test case before new system goes live — no untested rules in production
- Rule traceability: every rule in new system traced to either DER artifact, DOL approval, or incumbent documentation
- Legacy defect inventory: rules discovered to be legacy errors documented separately; DOL decides whether to perpetuate or correct each one
- Post-migration verification: 90-day post-cutover monitoring compares submission outcomes against legacy baseline to catch any missed rules
- Ongoing discovery: rules surfaced by filer complaints or data anomalies documented and added to rule registry during first year of operation

F5. Transition Support for External Integrators

Third-Party Transition Plan

- 90-day advance notice: all registered third-party software providers notified of API changes, schema updates, and cutover schedule
- Dedicated transition sandbox: production-mirrored environment available for third-party testing 90 days before cutover
- API version compatibility bridge: JSON and XML APIs supported concurrently during transition window; providers migrate on their own timeline
- Third-party recertification: all software providers must recertify via updated PATS before cutover date; certification available from day one of transition period
- Developer portal migration guide: step-by-step guide published for every API change with before/after examples and code snippets
- Dedicated integration support team: named Satsyl engineers available during transition period to assist third-party developers
- Drop-in office hours: weekly virtual sessions during transition period for third-party developers to get real-time guidance
- Annual developer conference: in-person and virtual conference each year to discuss upcoming changes; transition topics featured prominently

Cutover Support & Monitoring

- IRS/PBGC/DOL drop-box integrations: all downstream agency integrations tested and certified in staging before cutover — no surprises at go-live
- Third-party cutover monitoring: submission volumes and error rates monitored by software provider during cutover window
- Hotline during cutover: dedicated technical support hotline for third-party providers during cutover window; staffed 24/7 for 72 hours post-cutover
- Graceful degradation: if a third-party provider has not migrated by cutover, legacy-format submissions accepted in compatibility mode for 30 days
- Cutover status dashboard: real-time dashboard showing submission volumes by provider; third parties can self-monitor their cutover status
- Post-cutover follow-up: Satsyl contacts all third-party providers 2 weeks post-cutover to confirm successful migration and resolve any issues
- Lessons learned: post-cutover review with major third-party providers to capture feedback for future transition planning
- Ongoing support: dedicated third-party developer support channel maintained permanently — not just during transition

G. Third-Party and Bulk Filing

G1. High-Volume Submission Intake

Intake Architecture for Scale

- Amazon SQS intake queue: every bulk submission durably queued the moment it is received — intake API remains responsive even if processing is saturated
- Horizontal scaling: bulk API workers on Amazon EKS auto-scale; each worker handles thousands of submissions per hour
- API rate limiting: per-provider throttling protects system integrity while maximizing legitimate throughput; limits configurable by DOL
- Burst capacity: SQS queues absorb sudden deadline-day filing spikes without any degradation to the intake API response time
- Multi-format intake: JSON-first API with backward-compatible XML accepted in same endpoint; format detected automatically
- Batch submission: single API call accepts multiple filings; each filing within the batch independently processed and acknowledged
- Bulk file upload: SFTP endpoint available for providers who prefer file-based transfer; files automatically ingested into SQS pipeline
- Proven at scale: Satsyil's USPTO EDS platform processes 20M+ transactions/day — EFAST2's 1.4M annual filings is well within demonstrated capacity

Intake Monitoring & SLAs

- Real-time intake dashboard: queue depth, ingestion rate, and error rate visible to Satsyil ops and DOL in real time
- Intake SLA: acknowledgment issued within 30 seconds of receipt at peak load (95th percentile)
- Queue depth alerts: CloudWatch alarm triggers auto-scaling when queue depth exceeds threshold — prevents backlog buildup
- Provider-level monitoring: submission volumes tracked by software provider; anomalies flagged for investigation
- Intake health check: dedicated health endpoint for third parties to verify API availability before submitting batch runs
- Peak season preparation: intake capacity tested annually with JMeter/Locust simulating Columbus Day peak filing volume
- Graceful degradation: if downstream processing is impaired, intake continues accepting into durable queue; no submissions lost
- Incident notification: if intake API experiences degradation, registered third-party providers notified within 15 minutes via status page and email

G2. Processing Multiple Formats & Form Years

Multi-Format Processing

- RESTful JSON-first API: primary intake format for all new integrations; JSON schema auto-generated from current DER artifacts
- XML backward compatibility: prior-year XML submissions accepted via version-aware adapter; XML mapped to current internal data model automatically
- Format version routing: submission format and version detected on intake; routed to appropriate validation schema without third-party configuration
- Form type routing: Form 5500, 5500-SF, 5500-EZ, and all schedules processed through form-type-specific validation and business rule pipelines
- Mixed-batch handling: a single bulk submission may contain multiple form types and form years; each filing routed and processed independently
- Schema evolution: new DER versions introduced with namespace versioning; old versions continue processing until sunset date without system changes
- Format validation first: all submissions validated for well-formedness and schema compliance before any business rule processing
- Format error responses: format-level errors returned in same format

Multi-Year Coexistence

- Simultaneous form year processing: current year, prior year (amendments), and extension period submissions processed concurrently in production
- Form year tagging: every submission tagged with plan year on intake; validation engine applies form-year-specific rules automatically
- Amendment processing: amended filings linked to original submission; amendment history maintained and queryable for any plan
- Extended deadline handling: filings submitted under extension processed with appropriate form-year rules; extension status tracked and auditable
- Prior-year schema availability: validation schemas for all active form years maintained in production; no filer turned away for prior-year amendment
- Form year sunset: when a form year reaches end-of-life per DOL policy, submissions rejected with clear notification and contact information
- Multi-year reporting: DOL admin panel supports cross-year queries; year-over-year plan data comparison available for compliance analysis

G3. Secure Interfaces & APIs

API Security Architecture

- API gateway (Amazon API Gateway): all API traffic flows through managed gateway with WAF, rate limiting, authentication, and request validation
- API key authentication: third-party providers authenticate with scoped API keys issued after PATS certification; keys carry minimum necessary permissions
- NIST SP 800-63 key lifecycle: API keys created, distributed securely, rotated on schedule, and revoked immediately upon any compromise indication
- TLS 1.3 required: all API communications encrypted; older TLS versions and non-encrypted connections rejected at gateway
- FIPS 140-2 validated cryptography: all cryptographic operations use FIPS-validated modules — required for federal systems handling sensitive data
- Request signing: sensitive API operations require HMAC-SHA256 request signatures; prevents replay attacks and tampering
- IP allowlisting: third-party API keys optionally bound to registered IP ranges; requests from unexpected IPs flagged and may be blocked
- SFTP security: SSH key-pair authentication; password authentication disabled; keys rotated on credential lifecycle schedule

Developer Portal & Sandbox

- Developer portal: comprehensive documentation, API reference, code samples (Python/Java/C#), changelog, and support contact — always in sync with deployed API
- Sandbox environment: full-fidelity PATS sandbox available for testing; same validation engine, same APIs, same schemas as production
- Sandbox credential management: sandbox API keys issued separately from production; sandbox keys cannot access production data
- API versioning: semantic versioning with minimum 2-version support window; deprecation schedule published 12 months in advance
- API contract testing: automated monitoring verifies API responses match documented contracts; schema drift detected before reaching third parties
- Postman collection: maintained and published for all API endpoints; developers import and test immediately without manual setup
- API status page: real-time status for all API endpoints; historical uptime and incident history publicly visible to registered developers
- Security disclosure: responsible disclosure program for third parties who discover API security issues; issues triaged within 24 hours

G4. Acknowledgments, Status Updates & Error Reporting

Acknowledgment Flow

- Synchronous acknowledgment: bulk API returns acknowledgment within seconds of receiving each filing — before processing begins
- Unique tracking ID per filing: every filing in a bulk batch assigned independent tracking ID; providers track each filing individually
- Batch-level acknowledgment: batch submission acknowledged as a whole with individual filing references embedded
- Asynchronous status updates: processing status delivered via polling endpoint or push webhook as each filing progresses through the pipeline
- Webhook registration: third parties register callback URLs; status updates pushed within 30 seconds of each processing stage transition
- Email notifications: optional email alerts for submission accepted, errors found, and processing complete — configurable by the software provider
- Filer-level notifications: filers receive independent acknowledgment even when submission made through third-party software
- Final disposition notice: when filing fully distributed to DOL, IRS, and PBGC, final confirmation delivered with distribution timestamp

Error Reporting for Third Parties

- Structured error response: JSON error array with per-filing errors; each error includes filing reference, error code, field path, rule ID, and plain-language description
- Bulk error response: all errors for all filings in a batch returned in single response; third parties process errors programmatically
- Error severity included: each error classified as fatal (must fix to accept) or advisory (accepted with warning)
- Correction guidance: every error includes specific correction instruction — no cryptic codes without explanation
- Error deduplication: identical errors across multiple filings in a batch grouped; one fix corrects all affected filings
- Resubmission reference: rejected filings returned with reference ID used in resubmission; no need to re-send unchanged correct filings
- Error trend reporting: software providers can query their historical error rates by error code via API — supports self-service quality improvement
- Error escalation path: structured process for providers to escalate errors they believe are false positives; Satsyil responds within 2 business days

G5. Reliability Controls — Queuing, Retries & Exception Handling

Queue Management

- Amazon SQS decouples intake from processing: every submission durably stored in queue before processing begins — no data loss even if processing is temporarily unavailable
- Exactly-once delivery: SQS and AWS Step Functions use idempotency keys — each filing processed once and only once; no duplicates regardless of retry count
- Queue depth monitoring: CloudWatch tracks queue depth per provider and overall; auto-scaling triggers when depth exceeds threshold
- FIFO ordering: within each provider's submission batch, filings processed in submission order where order matters
- Dead-letter queue (DLQ): filings that exhaust all retries move to DLQ; ops team alerted within 60 seconds with filing ID, provider, and failure context
- DLQ monitoring: DLQ depth tracked separately; any DLQ entries require investigation and resolution — no silent discard
- Priority queuing: DOL can configure priority lanes for specific submission types or peak periods
- Queue visibility: Satsyil and DOL can view real-time queue state; queue metrics included in monthly operational report

Retry Logic & Exception Handling

- Automated retry with exponential backoff: transient failures retried automatically; backoff prevents thundering herd during recovery from outage
- Retry limits: configurable per error class; transient errors retry up to 5 times; permanent errors (validation failures) do not retry
- Retry visibility: each retry logged with attempt number, delay, and outcome; full retry history available for any filing
- Exception classification: transient (network/timeout), permanent (validation failure), or system (infrastructure) — different handling for each class
- System exception escalation: infrastructure-level exceptions trigger PagerDuty alert to ops team; resolved within SLA regardless of time of day
- Provider notification on permanent failure: when a filing exhausts retries, provider receives structured failure notice with all applicable error codes
- Exception trending: exception rates by type tracked weekly; increases trigger root cause investigation before they become incidents
- Recovery SLA: no filing left in DLQ for more than 24 hours without resolution or explicit DOL-approved exception

G6. Transition Strategy for Existing Third-Party Filers

Pre-Cutover Transition Program

- Transition announced 12 months before cutover: all registered software providers notified via developer portal, email, and annual developer conference
- Transition guide published 9 months before cutover: detailed migration documentation covering every API change, schema difference, and new requirement
- Compatibility bridge: legacy XML API maintained alongside new JSON API for full transition window — providers choose their own migration pace
- Dedicated transition sandbox: production-identical PATS environment available for testing 9 months before cutover; providers certify on their schedule
- Phased PATS recertification: certification open from 9 months before cutover; early-completing providers receive priority support at go-live
- Integration support team: named Satsyl engineers assigned to assist top-volume providers with integration questions during transition period
- Weekly office hours: virtual drop-in sessions for all third-party developers during transition period; recorded and posted for async viewing
- Annual developer conference: full-day event covering all transition

Cutover & Post-Cutover Support

- Cutover readiness dashboard: each provider's certification status visible; DOL and Satsyl monitor overall readiness before approving cutover date
- 72-hour cutover support: dedicated technical hotline staffed 24/7 for 72 hours post-cutover; major providers have named point of contact
- Compatibility mode: providers not certified by cutover accepted in legacy-format compatibility mode for 30 days; DOL approves extension policy
- Post-cutover monitoring: submission volumes and error rates tracked by provider for 90 days; providers with anomalies contacted proactively
- Post-cutover follow-up: Satsyl contacts all providers 2 weeks post-cutover to confirm successful migration and resolve any issues
- Legacy API sunset: after full transition window, legacy API deprecated with 90-day notice; no provider left without DOL approval
- Lessons learned: post-cutover retrospective with major providers and DOL to capture feedback and improve future transition processes
- Permanent support channel: dedicated third-party developer support channel maintained indefinitely — not just during transition window

H. Contact Center

H1. Accessibility for All Inquirers

Multi-Channel Accessibility

- AWS Connect cloud contact center: toll-free phone, email, web chat, and self-service portal — all channels available to filers and third-party providers
- Amazon Lex AI chatbot and IVR: 24/7 automated self-service for common inquiries (filing status, deadline information, password reset)
- Web portal self-service: authenticated filers access submission status, error details, and filing history without contacting an agent
- Email support: structured email intake with automated routing to appropriate specialist queue; response SLA tracked and reported
- Callback scheduling: filers who cannot wait on hold can schedule a callback at a preferred time; completed within scheduled window
- Multilingual support: translation services integrated for non-English speaking inquirers; primary languages identified from filer demographics
- All channels accessible from mobile device: responsive contact portal works on phone and tablet; no desktop-only limitations
- After-hours automated response: automated system acknowledges after-hours inquiries with expected response time and self-service options

Section 508 & Disability Access

- WCAG 2.2 AA compliance: all contact center web interfaces meet accessibility standards; tested quarterly with DHS Trusted Tester methodology
- TTY/TDD relay service: hearing-impaired callers supported via Federal Relay Service (1-800-877-8339); agents trained on relay call etiquette
- Screen reader compatibility: contact portal fully navigable by keyboard and screen reader; no mouse-only interactions
- Captioning: all recorded training and informational videos captioned; live captioning service available on request
- Cognitive accessibility: contact portal designed with plain language, clear navigation, and simple task flows per Plain Language Act
- Agent disability awareness training: all agents trained on accessible communication and accommodating diverse needs
- Accessibility testing: automated accessibility scanning on every portal deployment; manual review by accessibility tester annually
- Feedback mechanism: inquirers can report accessibility barriers through dedicated channel; issues triaged and resolved within 30 days

H2. Recording Inquiries & Contractor Responses

Call Recording Architecture

- AWS Connect native recording: all inbound and outbound calls recorded automatically; recording begins at call connection and ends at disconnect
- Screen recording: agent screen activity recorded during calls on request (configurable by DOL); supports dispute resolution and training
- Chat transcript recording: all web chat interactions logged as full transcripts with timestamps; archived alongside call recordings
- Email thread archiving: all email inquiries and responses archived in CRM (Salesforce/ServiceNow) with full thread history
- Voicemail recording: voicemails recorded and transcribed via Amazon Transcribe; text transcript attached to CRM case automatically
- Recording metadata: every recording tagged with call ID, agent ID, filer ID (if authenticated), inquiry category, and call duration
- Recording quality: stereo recording preserves agent and caller audio separately; enables independent quality review of each party
- Consent and notification: callers notified of recording per applicable regulations; consent documented in call record

Storage & Retrieval

- Secure storage: recordings stored in encrypted S3 bucket with KMS CMK; access restricted to authorized personnel only
- Retention policy: recordings retained for minimum 3 years per federal records requirements; configurable per DOL policy
- Search and retrieval: recordings searchable by call ID, agent, date range, inquiry category, and filer ID; retrieved in seconds
- Transcript search: Amazon Transcribe generates searchable text transcripts of all recordings; keyword search across all recorded calls
- DOL access: DOL authorized personnel retrieve any recording via secure portal without requesting through Satsyil
- Export capability: recordings exportable in standard formats (MP3/WAV) for DOL review or legal proceedings
- Audit log: every recording access logged with requester identity, access time, and purpose; tamper-evident
- Availability SLA: recordings available for retrieval within 1 hour of call completion; archived recordings retrieved within 24 hours

H3. Live & Recorded Inquiries for Monitoring

Government Monitoring Capabilities

- Live monitoring: DOL authorized personnel can listen to any active call in real time via AWS Connect supervisor dashboard
- Live dashboard: real-time view of queue depth, average wait time, active calls, and agent availability — DOL-accessible at all times
- Barge-in capability: DOL supervisor can join a live call to provide guidance to agent or speak directly with filer
- Silent monitoring: DOL monitors calls without agent or caller awareness — supports unbiased quality assessment
- Recording review portal: DOL staff access recorded calls via secure web portal; filter by date, agent, inquiry type, or filer ID
- Quality scorecard review: DOL can view Satsyil quality scores for any call; scores based on standardized rubric agreed with DOL
- Trend reports: DOL receives monthly reports on call volume, handle time, first-call resolution rate, and quality scores
- Ad-hoc query: DOL can request call samples meeting specific criteria (e.g., calls about a specific form type) for targeted review

Contractor Quality Monitoring

- Contractor QA team: dedicated QA analysts review minimum 5% of calls per agent per week using standardized quality rubric
- AI quality scoring: Amazon Connect Contact Lens analyzes all calls for sentiment, dead air, compliance phrases, and script adherence
- Quality rubric: scoring criteria developed jointly with DOL; covers accuracy, professionalism, accessibility, and issue resolution
- Agent feedback: QA results shared with agents within 48 hours; coaching sessions scheduled for scores below threshold
- Agent training: recordings used as training material; best-practice calls shared with all agents; problematic patterns addressed in team training
- Performance improvement plans: agents with sustained quality issues enrolled in structured improvement program with documented milestones
- Quality trend monitoring: if quality scores drop across multiple agents, systemic issue investigation triggered; DOL notified of findings
- Monthly quality report to DOL: overall quality scores, distribution by rubric category, top issues, and improvement actions taken

H4. Cost-Effective Solutions Without High Licensing Costs

AWS Connect Cost Model

- Pay-per-minute pricing: AWS Connect charges per active call minute — no per-seat licensing fees; estimated 40-60% cost reduction vs. legacy or high-license solutions
- No hardware procurement: AWS Connect is fully cloud-based; zero on-premise hardware, PBX, or telephony equipment required
- Existing DOL toll-free number ported to AWS Connect: no new phone number procurement; seamless transition for filers
- Browser-based softphones: agents work from any device with a browser and headset — no proprietary desk phones, no VPN required
- Elastic capacity: agent seats activated and deactivated dynamically; no minimum seat commitment; scale up for October peak, scale down in off-season
- Amazon Lex chatbot: AI self-service handles 40%+ of routine inquiries (status checks, deadline questions) at a fraction of agent cost
- IVR deflection: well-designed IVR resolves simple inquiries before reaching an agent; deflection rate tracked and optimized continuously
- No additional CRM licensing: AWS Connect integrates with ServiceNow or Salesforce (DOL's existing systems) at no additional CRM cost

Cost Transparency & Governance

- Monthly cost reporting: AWS Connect usage costs reported to DOL monthly with breakdown by channel (phone, chat, IVR) and agent hour
- Cost per contact metric: DOL tracks cost-per-inquiry monthly; target baseline established in contract; continuous improvement tracked against target
- Workforce management: predictive staffing model scales agents to match forecast volume; prevents overstaffing during off-peak periods
- Peak season staffing: contractor surges agent count for October and January peak periods using cross-trained staff; no separate licensing cost
- Self-service investment: investment in IVR and chatbot improvements reduces live agent volume over time; ROI reported to DOL annually
- Cost avoidance reporting: DOL receives annual report estimating cost avoidance from IVR deflection, chatbot handling, and first-call resolution
- Vendor comparison: if DOL requests, Satsyl provides annual comparison of AWS Connect costs vs. alternative solutions; no hidden costs
- Exit-ready: no proprietary contact center lock-in; call recordings,

Similar Application Demonstration

Similar Application Demonstration

1. Dynamic Form Entry & Validation: Live submission workflow with field-level validation, cross-field business rule enforcement, and real-time error messaging with correction guidance.


2. Attachment & Document Handling: Upload, virus scan, and secure attachment to a submission record; PII detection and automated redaction demonstrated live.

3. Submission Lifecycle & Status Tracking: End-to-end flow: intake queue → validation → processing → acknowledgment → status dashboard with unique tracking ID.

4. Admin & Configuration Panel: Adding/modifying a business rule without code; form year version management and schema update, no deployment required.

5. Bulk Filing API: Third-party bulk submission via API with structured error response, retry handling, and PATS certification workflow.

6. Contact Center Integration: CRM inquiry log tied to a specific submission record; call recording access and agent knowledge base demonstrated.



Similar Application: USPTO
Enterprise Data Services (EDS)
Platform — 20M+ transactions/day,
10+ PB federal data

Fallback Video: 5-minute pre-recorded
video prepared covering all 6 demo
areas above — available immediately
if any live technical issue arises

Government Areas of Interest

Key Interests

How Satsyil Addresses DOL's Key Areas of Interest

- **Architecture at Scale:** Multi-AZ EKS + SQS queuing handles 1.4M+ filings; CDN serves 90% of static traffic; auto-scaling fires at 70% CPU — zero manual intervention.
- **Quality Controls & Queuing:** Every submission enters SQS for exactly-once processing; dead-letter queues capture failures; automated retry with exponential backoff; ops team alerted within 60 seconds.
- **Annual Changes:** Config-driven rules engine, DOL staff modify fields/rules in admin panel, no code deployment, no downtime. Third parties get 90-day advance notice; backward-compatible APIs maintained.
- **System Currency & Security:** Continuous cATO via DevSecOps; SAST/DAST on every commit; automated STIG compliance; GuardDuty + Security Hub real-time detection; critical CVE patching within 15 days.
- **Migration Validation:** Parallel run with field-level reconciliation at every ETL stage; automated row counts, checksums, and business rule re-execution on migrated data; phased go/no-go cutover gates.
- **External Integrator Transition:** 90-day advance notice; dedicated sandbox; API version bridge (JSON + XML); third-party recertification via updated PATS; dedicated integration support team at cutover.
- **Contact Center:** AWS Connect with Lex IVR for 24/7 self-service; all calls recorded with searchable transcripts; government has live dashboard access and call recording review rights; quality scored automatically.
- **AI for Delivery:** AI-assisted code gen and test case generation accelerate development; ML models detect PII/FTI pre-release; anomaly detection flags irregular patterns; all AI governed with human-in-the-loop for consequential actions.

Q&A

Questions & Answers

Venugopal Ankinapalli, President

Email: venugopal.ankinapalli@satsyil.com

Tel: (703) 689-1842 | www.satsyil.com